

# AI Project

QuestCrafter: Build an AI “Dungeon Master” for RPG Quests

MSc1 — AI

# Project Overview

**QuestCrafter** is a small text-generation system that creates **short fantasy quests** from a user prompt, e.g.:

- "Create a level-3 quest in a desert city with a betrayal twist."

**Main idea:** start from an existing small language model, build a baseline, then fine-tune it with a curated dataset, and deliver a simple demo.

**You will practice:**

- Data preprocessing for instruction-style generation (prompt → completion)
- Baseline vs fine-tuned model comparison
- Lightweight training with **PyTorch** (via **Hugging Face** transformers)
- Evaluation (automatic + human rubric)
- Demo app (Streamlit or Gradio)

# Learning Objectives

By the end of week 5, you should be able to:

- ① Prepare a text dataset in a clean supervised format (JSONL/CSV).
- ② Run inference with a pretrained model and analyze typical failures.
- ③ Fine-tune a small model with **PyTorch** (optionally with LoRA/PEFT).
- ④ Evaluate generations using a fixed test set + a human scoring rubric.
- ⑤ Deliver a reproducible mini-product: repo + report + demo.

# Dataset (Choose One)

Pick **one** dataset for training:

- **WritingPrompts (Kaggle)**: prompt/story pairs (excellent for controlled prompting).
- **TinyStories (Hugging Face)**: short, clean stories (easy and fast to train/evaluate).
- **Reddit Jokes (Kaggle)** (optional alternative): style control for humor.

**Recommended for beginners:** WritingPrompts or TinyStories.

**Target training format:**

- instruction/prompt → response (quest text)
- Optional metadata: level, setting, tone, length

# Modeling Approach (with PyTorch)

**Framework:** PyTorch + Hugging Face transformers (PyTorch backend)

**Suggested small models (CPU/GPU friendly):**

- Decoder-only: distilgpt2 or gpt2
- Encoder-decoder: t5-small (instruction-friendly)

**Training options:**

- **Full fine-tuning** (simple, but heavier)
- **LoRA / PEFT** (recommended: faster, cheaper, easy to compare)

**Implementation tip:** start with distilgpt2 for quick results.

# End-to-End Pipeline

- ① **Data** download → cleaning → formatting (JSONL).
- ② **Baseline inference:** generate quests with pretrained model (no training).
- ③ **Fine-tuning:** train on your prompt → quest pairs.
- ④ **Control:** add fields like LEVEL, SETTING, TONE in prompts.
- ⑤ **Evaluation:** automatic metrics + human rubric + example comparisons.
- ⑥ **Demo:** small app for interactive generation.

# Timetable (5 Weeks)

Week	Main activities and outputs
W1	Dataset selection + exploration; cleaning and formatting; define quest template fields; build train/val/test splits.
W2	Baseline generation (pretrained model); create fixed test prompts (50+); define evaluation rubric; first failure analysis.
W3	Fine-tune one small model in <b>PyTorch</b> (optionally LoRA); track training curves; save checkpoints; sample outputs after training.
W4	Add controllability (level/setting/tone/length); add safety filters; structured evaluation: baseline vs tuned; robustness tests.
W5	Build demo (Streamlit/Gradio); finalize report (4–8 pages); create final repo (scripts + README); final presentation.

# Timetable Matrix (What Happens When)

Task	W1	W2	W3	W4	W5
Dataset pipeline (clean + format)	✓				
Baseline inference + analysis		✓			
Fine-tuning (PyTorch)			✓		
Control + safety + robustness				✓	
Evaluation + reporting	✓	✓	✓	✓	
Demo app + packaging					✓

Note: ✓ indicates the main focus of the week (some tasks continue across weeks).

# # Mandatory Tasks (Required)

## ① Dataset pipeline

- Download dataset, clean, format to JSONL/CSV: prompt → response
- Split into train/val/test (e.g., 80/10/10)

## ② Baseline generation (no training)

- Generate outputs using a pretrained model
- Document typical issues (repetition, off-topic, too long, incoherence)

## ③ Fine-tuning with PyTorch

- Train at least one model (distilgpt2/gpt2/t5-small)
- Provide config (epochs, LR, batch size, max length)

## ④ Evaluation

- Fixed test set of **50+ prompts**
- Automatic metric: validation loss / perplexity + one diversity metric (e.g., Distinct-n)
- Human rubric: coherence, creativity, prompt-faithfulness (1–5)

## ⑤ Final deliverable

- Repo + README + report (4–8 pages) + demo app

## # Optional Tasks (Choose Any)

- **LoRA/PEFT** fine-tuning and compare vs full fine-tuning
- Add **control tokens** like <LEVEL=3><TONE=dark><SETTING=forest>
- **RAG add-on:** retrieve from a small “lore book” before generation
- Multi-output: quest + NPC dialogue + item description + reward summary
- Automatic quality filters (repetition detection, length control, regeneration)
- Class leaderboard: compare models with the same test prompts + rubric
- Write a short **model card** (limitations, bias, safety, intended use)

# Evaluation (Baseline vs Fine-Tuned)

**Fixed test prompts:** 50+ prompts covering:

- different settings (forest, desert, cyberpunk, medieval city)
- different levels (1–10) and lengths (short/medium)
- different tones (epic, humorous, dark)

**Human rubric (1–5 each):**

- Coherence (logical story flow)
- Prompt-faithfulness (respects constraints)
- Creativity (interesting elements, not generic)

**Automatic:** validation loss/perplexity + Distinct-n (diversity).

# Final Deliverables

## 1) Git repository

- data/ (or script to download + preprocess)
- preprocess.py (or notebook)
- train.py (PyTorch training via **Hugging Face Trainer** or custom loop)
- evaluate.py (metrics + saving generations)
- demo\_app.py (*Streamlit/Gradio*)
- README.md (how to run + results summary)

## 2) Report (4–8 pages)

- Dataset + formatting, baseline, fine-tuning setup, evaluation protocol, results

## 3) Presentation

- demo + baseline vs tuned examples + what you learned

# Suggested Grading Breakdown (Example)

Criterion	Weight
Dataset pipeline quality + reproducibility	20%
Baseline + analysis (failure cases, examples)	15%
Fine-tuning (correctness, experiment tracking)	25%
Evaluation protocol (test set + rubric + metrics)	20%
Demo + final packaging (repo + README + usability)	20%

# Practical Notes (Compute & Tips)

- Start small: use a subset of the dataset for quick iteration.
- Keep generations short at first (e.g., max 120–200 tokens).
- Track everything: seed, hyperparameters, checkpoints, prompt set.
- Prefer **LoRA/PEFT** if GPU memory is limited.
- Always compare against baseline with the **same** test prompts.

## Common failure modes to monitor:

- repetition loops, off-topic drift, ignored constraints, over-long outputs

# Conclusion

**QuestCrafter** is a beginner-friendly, end-to-end generative AI project:

- real dataset + real training in **PyTorch**
- clear mandatory milestones
- optional extensions (LoRA, RAG, control tokens)
- a demo you can show in interviews

*Next step: pick the dataset (WritingPrompts or TinyStories) and build W1 pipeline.*